

# HRL1\_final

April 23, 2022

```
[ ]: import numpy as np
import random
import gym
from gym.wrappers import Monitor
import glob
import io
import matplotlib.pyplot as plt
from IPython.display import HTML
import pandas as pd
from IPython import display
from IPython.display import clear_output
from tqdm import tqdm
import time
from IPython.display import display, clear_output
from mpl_toolkits.axes_grid1 import make_axes_locatable
import matplotlib.patches as patches
```

```
[ ]: #Setting up the environment
env = gym.make('Taxi-v3')

env.reset()

#Current State
print(env.s)

print ("Number of states:", env.nS)

# Primitive Actions
action = ["south", "north", "east", "west", "pick", "drop"]
#correspond to [0,1,2,3] that's actually passed to the environment

# R, G, Y, B, 0 , 1, 2, 3 in state decoding
# either go left, up, down or right, pickup or dropoff
print ("Number of actions that an agent can take:", env.nA)

# Example Transitions
rnd_action = random.randint(0, 5)
```

```

print ("Action taken:", action[rnd_action])
next_state, reward, is_terminal, t_prob = env.step(rnd_action)
print ("Transition probability:", t_prob)
print ("Next state:", next_state)
print ("Reward recieved:", reward)
print ("Terminal state:", is_terminal)
env.render()

```

269

Number of states: 500

Number of actions that an agent can take: 6

Action taken: west

Transition probability: {'prob': 1.0}

Next state: 249

Reward recieved: -1

Terminal state: False

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(West)

```

```

[ ]: env.reset()
env.step(0)

env.reset()
# print(list(env.decode(env.s))[:2])
print(list(env.decode(env.s)))
env.render()
R = [0,0]
G = [0,4]
Y = [4,0]
B = [4,3]

```

```

[0, 4, 2, 1]
+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+

```

We consider the goal state to be either of the four R,G,Y,B. We need a data structure to hold these

4 Q-tables.

```
[ ]: seed = 42
     rg = np.random.RandomState(seed)

[ ]: def plot_Q(Q, message = "Q plot"):

    # plt.figure(figsize=(10,10))
    fig, ax = plt.subplots(figsize = (10,10))

    Q_max = Q.max(-1)

    im = ax.imshow(Q_max)

    cbar = ax.figure.colorbar(im, ax=ax)

    ax.set_title(message)
    #ax.grid(visible = True, which="major", color="w", linestyle='-', linewidth=2)
    ax.set_xlim([0, 5])
    ax.set_ylim([0,5])
    ax.pcolor(Q_max, edgecolors='k', linewidths=1)
    #plt.colorbar()

    # plt.colorbar()
    def x_direct(a):
        if a in [0, 1]:
            return 0
        return 1 if a == 2 else -1
    def y_direct(a):
        if a in [2, 3]:
            return 0
        return 1 if a == 1 else -1
    policy = Q.argmax(-1)
    policyx = np.vectorize(x_direct)(policy)
    policyy = np.vectorize(y_direct)(policy)
    idx = np.indices(policy.shape)
    ax.quiver(idx[1].ravel()+ 0.5, idx[0].ravel()+0.5, policyx.ravel(), policyy.
    ravel(), pivot="middle", color='red')
    fig.tight_layout()

def q_learning(env,goal_state, episodes = 1000, alpha0 = 0.4, epsilon0 = 0.2,
    beta = 1, gamma = 0.9, plot_heat = True, print_freq = 100, max_steps = 200):
```

```

Q = np.zeros((5,5, 4)) #only first 4 actions allowed

episode_rewards = np.zeros(episodes)
steps_to_completion = np.zeros(episodes)
'''if plot_heat:
    clear_output(wait=True)
    plot_Q(Q)'''
epsilon = epsilon0
alpha = alpha0
for ep in tqdm(range(episodes)):
    tot_reward, steps = 0, 0

    # Reset environment
    state_seq = env.reset()
    state = list(env.decode(state_seq))[:2]

    done = False
    while not done:
        action = egreedy_policy(Q,state,epsilon)

        state_next_seq, reward ,_,_ = env.step(action)

        #action_next = np.argmax(Q[state_next])
        state_next = list(env.decode(state_next_seq))[:2]

        if(state_next == state):
            reward += -20

        if(state_next == goal_state):
            reward += 100

        # update equation
        Q[state[0],state[1], action] += alpha*(reward + gamma*np.
→max(Q[state_next[0],state_next[1]]) - Q[state[0],state[1], action])

        tot_reward += reward
        steps += 1

        #How do we know the state index of R,G,Y,B?
        if steps == max_steps or state_next == goal_state : done =True

        state = state_next

    episode_rewards[ep] = tot_reward
    steps_to_completion[ep] = steps

```

```

        #if (ep+1)%print_freq == 0 and plot_heat:
        clear_output(wait=True)
        plot_Q(Q, message = "Episode %d: Reward: %f, Steps: %.2f, Qmax: %.2f, Qmin:
↪%.2f"%(ep+1, np.mean(episode_rewards[ep-print_freq+1:ep]),
np.
↪mean(steps_to_completion[ep-print_freq+1:ep]),
Q.
↪max(), Q.min()))
        #display.display( df.T )
        return Q, episode_rewards, steps_to_completion, env

```

```

[ ]: from IPython.display import Image
# alpha = 0.4
alpha= 0.19751633358692788
epsilon = 0.12233772892858322
gamma = 0.9684975038393936

R_mat = np.array([[1,3,0,0,0],
                  [1,1,0,0,0],
                  [1,1,3,3,3],
                  [1,1,1,1,1],
                  [1,1,1,1,1]])
#Image('/content/option_R.png')

```

```

[ ]: Y_mat = np.array([[0,0,0,0,0],
                      [0,0,0,0,0],
                      [0,3,3,3,3],
                      [0,1,1,1,1],
                      [0,1,1,1,1]])
#Image('/content/option_Y.png')

```

```

[ ]: G_mat = np.array([[0,0,2,2,1],
                      [0,0,1,1,1],
                      [2,2,2,2,1],
                      [1,1,1,1,1],
                      [1,1,1,1,1]])
#Image('/content/option_G.png')

```

```

[ ]: B_mat = np.array([[0,0,0,0,0],
                      [0,0,0,0,0],
                      [2,2,2,0,0],
                      [1,1,1,0,0],
                      [1,1,1,0,3]])
#Image('/content/option_B.png')

```

```

[ ]: # We are defining four more options here
# Option 1 move_to_R
# Option 2 move_to_G
# Option 3 move_to_Y
# Option 4 move_to_B

def move_to_R(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = R_mat[ds[0],ds[1]]

    if (ds == [0,0]):    #termination condition
        optdone = True

    return [optact,optdone]

def move_to_Y(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = Y_mat[ds[0],ds[1]]

    if (ds == [4,0]):    #termination condition
        optdone = True

    return [optact,optdone]

def move_to_G(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = G_mat[ds[0],ds[1]]

    if (ds == [0,4]):    #termination condition
        optdone = True

    return [optact,optdone]

def move_to_B(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = B_mat[ds[0],ds[1]]

    if (ds == [4,3]):    #termination condition
        optdone = True

```

```
return [optact,optdone]
```

```
#Now the new action space will contain  
#Primitive Actions: ["south", "north", "east", "west", "pick", "drop"]  
#Options: ["move_to_R", "move_to_Y", "move_to_G", "move_to_B"]  
#Total Actions :["south", "north", "east",  
↪ "west", "pick", "drop", "move_to_R", "move_to_Y", "move_to_G", "move_to_B"]  
#Corresponding to [0,1,2,3,4,5,6,7,8,9]
```

```
[ ]: actions=[0,1,2,3,4,5,6,7,8,9]  
#epsilon-greedy action selection function  
seed = 36  
rg = np.random.RandomState(seed)  
  
def egreedy_policy(q_values,state,epsilon):  
    if rg.rand() < epsilon:  
        return rg.choice(actions)  
    else:  
        #max = np.max(q_values[state])  
        #return rg.choice(np.where(q_values[state] == max)[0])  
        return np.argmax(q_values[state])
```

```
[ ]: # sweep_config = {  
#     "name" : "SMDP-sweep",  
#     "method": "random",  
#     "parameters": {  
#         "gamma": {  
#             "min": 0.850,  
#             "max": 0.999  
#         },  
#         "alpha": {  
#             "min": 0.1,  
#             "max": 0.4  
#         },  
#         "epsilon": {  
#             "min": 0.1,  
#             "max": 0.3  
#         }  
#     }  
# }
```

```
[ ]: #sweep_id = wandb.sweep(sweep_config, project='RLPA3')
```

```
[ ]: #wandb.agent(sweep_id, SMDP, count=20)
```

```

[ ]: ##### SMDP Q-Learning

# Add parameters you might need here
# gamma = 0.9
# alpha = 0.4
alpha1 = 0.2821687004024015
epsilon1 = 0.13732834786589931
gamma1 = 0.9756392557916136
def SMDP(gamma, alpha, epsilon):
    # wandb.init(project = 'RLPA3', entity = 'reinforce-boys')
    # gamma = wandb.config.gamma
    # alpha = wandb.config.alpha
    # epsilon = wandb.config.epsilon

    # Iterate over 1000 episodes
    q_values_SMDP = np.zeros((500,10))
    ufd1 = np.zeros((500,10)) #Update_Frequency Data structure
    Rewards = []
    for _ in range(50000):
        state = env.reset()
        done = False

        # While episode is not over
        episode_reward = 0
        while not done:

            # Choose action
            action = egreedy_policy(q_values_SMDP, state, epsilon=epsilon)

            # Checking if primitive action
            if action < 6:
                # Perform regular Q-Learning update for state-action pair

                next_state, reward, done, _ = env.step(action)
                q_values_SMDP[state, action] += alpha*(reward + gamma*np.
→max([q_values_SMDP[next_state, action] for action in actions]) -
→q_values_SMDP[state, action])
                ufd1[state,action] += 1
                state = next_state
                episode_reward += reward

            # Checking if action chosen is an option
            reward_bar = 0
            if action == 6: # action => move_to_R

                initial_state = np.copy(state)
                optdone = False

```



```

count=0
while (optdone == False):

    optact,_ = move_to_R(env,state)
    #
    next_state, reward, done,_ = env.step(optact)

    _,optdone = move_to_R(env,next_state)

    reward_bar = reward_bar + (gamma**count)*reward
    count+=1
    state = next_state
    episode_reward += reward

    q_values_SMDP[initial_state, action] += alpha*(reward_bar +
→(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
→- q_values_SMDP[initial_state, action])
    ufd1[initial_state,action] += 1

if action == 7: # action => move_to_Y option

    initial_state = np.copy(state)
    optdone = False
    count=0
    while (optdone == False):

        optact,_ = move_to_Y(env,state)

        next_state, reward, done,_ = env.step(optact)

        _,optdone = move_to_Y(env,next_state)

        reward_bar = reward_bar + (gamma**count)*reward
        count+=1

        state = next_state
        episode_reward += reward

        q_values_SMDP[initial_state, action] += alpha*(reward_bar +
→(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
→- q_values_SMDP[initial_state, action])
        ufd1[initial_state,action] += 1

if action == 8: # action => move_to_G option

```

```

initial_state = np.copy(state)
optdone = False
count=0
while (optdone == False):

    optact,_ = move_to_G(env,state)
    #
    next_state, reward, done,_ = env.step(optact)

    _,optdone = move_to_G(env,next_state)

    reward_bar = reward_bar + (gamma**count)*reward
    count+=1

    state = next_state
    episode_reward += reward

    q_values_SMDP[initial_state, action] += alpha*(reward_bar +
→(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
→- q_values_SMDP[initial_state, action])
    ufd1[initial_state,action] += 1

if action == 9: # action => move_to_B option

    initial_state = np.copy(state)
    optdone = False
    count=0
    while (optdone == False):

        optact,_ = move_to_B(env,state)
        #
        next_state, reward, done,_ = env.step(optact)

        _,optdone = move_to_B(env,next_state)

        reward_bar = reward_bar + (gamma**count)*reward
        count+=1

        state = next_state
        episode_reward += reward

        q_values_SMDP[initial_state, action] += alpha*(reward_bar +
→(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
→- q_values_SMDP[initial_state, action])

```

```

        ufd1[initial_state,action] += 1

        #wandb.log({'episode_reward': episode_reward})
        Rewards.append(episode_reward)
    return q_values_SMDP, Rewards, ufd1

```

```

#plt.plot(Rewards)

```

```

[ ]: q_values_SMDP, Rewards, ufd1 = SMDP(gamma1, alpha1, epsilon1)

```

```

[ ]: np.mean(Rewards[40000:])

```

```

[ ]: 4.0599

```

```

[ ]: pass_loc = ['R','G','Y','B','in_taxi']
def render_taxi(s):
    env.s = s
    #plt.figure(3)
    #plt.clf()
    #plt.imshow(env.render(mode='rgb_array'))
    env.render()
    #display.clear_output(wait=True)
    #display.display(plt.gcf())
    #env.render()
    time.sleep(2)
    clear_output(wait = True)

def drive_taxi(q_values):
    state = env.reset()
    done = False
    env.render()
    clear_output(wait = True)
    print("Passenger location : "+pass_loc[list(env.decode(env.s))[2]])
    print("Destination : "+pass_loc[list(env.decode(env.s))[3]])
    while not done:

        action = np.argmax([q_values[state]])
        if action < 6:
            state, reward, done, _ = env.step(action)
            render_taxi(state)
        if action == 6:
            optdone = False
            while(optdone == False):
                optact, optdone = move_to_R(env,state)
                next_state, reward, done, _ = env.step(optact)

```

```

        state = next_state
        render_taxi(state)
    if action == 7:
        optdone = False
        while(optdone == False):
            optact, optdone = move_to_Y(env,state)
            next_state, reward, done,_ = env.step(optact)
            state = next_state
            render_taxi(state)
    if action == 8:
        optdone = False
        while(optdone == False):
            optact, optdone = move_to_G(env,state)
            next_state, reward, done,_ = env.step(optact)
            state = next_state
            render_taxi(state)
    if action == 9:
        optdone = False
        while(optdone == False):
            optact, optdone = move_to_B(env,state)
            next_state, reward, done,_ = env.step(optact)
            state = next_state
            render_taxi(state)

```

```

[ ]: #img = plt.imshow(env.render(mode = 'rgb_array'))
      drive_taxi(q_values_SMDP)

```

```

+-----+
|R: | : :G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Dropoff)

```

```

[ ]: # sweep_config = {
#     "name" : "Intra-sweep",
#     "method": "random",
#     "parameters": {
#         "gamma": {
#             "min": 0.850,
#             "max": 0.999
#         },
#         "alpha": {

```

```

#         "min": 0.1,
#         "max": 0.4
#     },
#     "epsilon": {
#         "min": 0.1,
#         "max": 0.3
#     }
# }
# }

```

```
[ ]: #sweep_id = wandb.sweep(sweep_config, project='RLPA3')
```

```
[ ]: #wandb.agent(sweep_id, SMDP, count=20)
```

```
[ ]: ##### Intra-Option Q-Learning
```

```

# Add parameters you might need here
# gamma = 0.9
# alpha = 0.4
alpha= 0.19751633358692788
epsilon = 0.12233772892858322
gamma = 0.9684975038393936
def Intra(gamma, alpha, epsilon):
    # Iterate over 1000 episodes
    # gamma = wandb.config.gamma
    # alpha = wandb.config.alpha
    # epsilon = wandb.config.epsilon
    Rewards = []
    q_values_SMDP2 = np.zeros((500,10))

    ufd2 = np.zeros((500,10))#Update_Frequency Data structure
    for _ in range(50000):
        state = env.reset()
        done = False

        # While episode is not over
        episode_reward = 0
        while not done:

            # Choose action
            action = egreedy_policy(q_values_SMDP2, state, epsilon=epsilon)

            # Checking if primitive action
            if action < 6:
                # Perform regular Q-Learning update for state-action pair

                next_state, reward, done, _ = env.step(action)

```

```

        q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) -
↪q_values_SMDP2[state, action])
        ufd2[state,action] += 1
        episode_reward+=reward

        state = next_state

# Checking if action chosen is an option

    if action == 6: # action => Move to R option

        optdone = False
        while (optdone == False) :

            optact,_ = move_to_R(env,state)
            next_state, reward, done,_ = env.step(optact)
            _,optdone = move_to_R(env,next_state)
            episode_reward+=reward

            q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) -
↪q_values_SMDP2[state, optact])
            ufd2[state,optact] += 1

            if not optdone:
                q_values_SMDP2[state, action] += alpha*(reward +
↪gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
                ufd2[state,action] += 1
            else:
                q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) -
↪q_values_SMDP2[state, action])
                ufd2[state,action] += 1

            state = next_state

    if action == 7: # action => move to Y option

        optdone = False
        while (optdone == False) :

            optact,_ = move_to_Y(env,state)

```

```

        next_state, reward, done, _ = env.step(optact)
        _, optdone = move_to_Y(env, next_state)
        episode_reward += reward

        q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) - □
↪q_values_SMDP2[state, optact])
        ufd2[state, optact] += 1

        if not optdone:
            q_values_SMDP2[state, action] += alpha*(reward + □
↪gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
            ufd2[state, action] += 1
        else:
            q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) - □
↪q_values_SMDP2[state, action])
            ufd2[state, action] += 1

        state = next_state

    if action == 8: # action => move to G option

        optdone = False
        while (optdone == False) :

            optact, _ = move_to_Y(env, state)
            next_state, reward, done, _ = env.step(optact)
            _, optdone = move_to_Y(env, next_state)
            episode_reward += reward

            q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) - □
↪q_values_SMDP2[state, optact])
            ufd2[state, optact] += 1

            if not optdone:
                q_values_SMDP2[state, action] += alpha*(reward + □
↪gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
                ufd2[state, action] += 1
            else:
                q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions])) - □
↪q_values_SMDP2[state, action])
                ufd2[state, action] += 1

```

```

        state = next_state

    if action == 9: # action => move to B option

        optdone = False
        while (optdone == False) :

            optact, _ = move_to_B(env, state)
            next_state, reward, done, _ = env.step(optact)
            _, optdone = move_to_B(env, next_state)
            episode_reward += reward

            q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
→max([q_values_SMDP2[next_state, action] for action in actions]) -
→q_values_SMDP2[state, optact])
            ufd2[state, optact] += 1

            if not optdone:
                q_values_SMDP2[state, action] += alpha*(reward +
→gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
                ufd2[state, action] += 1
            else:
                q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
→max([q_values_SMDP2[next_state, action] for action in actions]) -
→q_values_SMDP2[state, action])
                ufd2[state, action] += 1

            state = next_state
            #wandb.log({'episode_reward': episode_reward})
            Rewards.append(episode_reward)
        return q_values_SMDP2, Rewards, ufd2

# plt.plot(Rewards)

```

```
[ ]: q_values_SMDP2, Rewards2, ufd2 = Intra(gamma, alpha, epsilon)
```

```
[ ]: np.mean(Rewards2[40000:])
```

```
[ ]: -1.2759
```

```
[ ]: drive_taxi(q_values_SMDP2)
```

```

+-----+
|R: | : :G|

```



```

| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
(Dropoff)

```

```

[ ]: #case when passenger is at either R,G,Y,B
#To visualise lets break the task into 2 parts: pick-up and drop
op_SMDP_1 = [np.zeros((5,5)) for i in range(4)] #Matrix to hold the optimal_
    ↪actions/options in each state
op_SMDP_q1 = [np.zeros((5,5)) for i in range(4)]
op_intra_1 = [np.zeros((5,5)) for i in range(4)] #cases where passengers are_
    ↪not in the taxi
op_intra_q1 = [np.zeros((5,5)) for i in range(4)]

op_SMDP_2 = [np.zeros((5,5)) for i in range(4)] #Matrix to hold the optimal_
    ↪actions/options in each state
op_SMDP_q2 = [np.zeros((5,5)) for i in range(4)]
op_intra_2 = [np.zeros((5,5)) for i in range(4)] #cases where passenger is in_
    ↪the taxi
op_intra_q2 = [np.zeros((5,5)) for i in range(4)]

for i in range(500): #iterating over all states
    state = list(env.decode(i))
    if state[2] != 4:
        op_SMDP_1[state[2]][state[0],state[1]] = np.argmax(q_values_SMDP[i])
        op_SMDP_q1[state[2]][state[0],state[1]] = np.amax(q_values_SMDP[i])
        op_intra_1[state[2]][state[0],state[1]] = np.argmax(q_values_SMDP2[i])
        op_intra_q1[state[2]][state[0],state[1]] = np.amax(q_values_SMDP2[i])
    else:
        op_SMDP_2[state[3]][state[0],state[1]] = np.argmax(q_values_SMDP[i])
        op_SMDP_q2[state[3]][state[0],state[1]] = np.amax(q_values_SMDP[i])
        op_intra_2[state[3]][state[0],state[1]] = np.argmax(q_values_SMDP2[i])
        op_intra_q2[state[3]][state[0],state[1]] = np.amax(q_values_SMDP2[i])

```

```

[ ]: def visualise_q(m1,m2):
    fig, ax = plt.subplots(figsize = (10,10))
    im = ax.imshow(m1, extent=[0, 10, 0, 10])
    ax.grid(which='major', axis='both', linestyle='-', color='k', linewidth=2)
    divider = make_axes_locatable(ax)
    cax = divider.append_axes('right', size='5%', pad=0.05)
    fig.colorbar(im, cax=cax, orientation='vertical')
    def x_direct(a):
        if a in [0,1,4,5]:
            return 0
        elif a in [2,8,9]:

```

```

        return 1
    else:
        return -1
def y_direct(a):
    if a in [2,3,4,5]:
        return 0
    elif a in [1,6,8]:
        return 1
    else:
        return -1
policyx = np.vectorize(x_direct)(m2)
polycyy = np.vectorize(y_direct)(m2)
idx = 2*np.indices((5,5))
ax.quiver(idx[1][::-1].ravel()+ 1, idx[0][::-1].ravel()+1, policyx.ravel(),
↪polycyy.ravel(), pivot="middle", color='red')

for i in range(5):
    for j in range(5):
        if m2[i][j] == 4: #pick-rectangle
            rect = patches.Rectangle((idx[1][::-1][i][j]+0.5, idx[0][::-1][i][j]+0.
↪5), 1, 1, linewidth=0.5, edgecolor='r', facecolor='r')
            ax.add_patch(rect)
        if m2[i][j] == 5: #drop-Circle
            circ = patches.Circle((idx[1][::-1][i][j]+1, idx[0][::-1][i][j]+1), 0.
↪5, linewidth=0.5, edgecolor='r', facecolor='r')
            ax.add_patch(circ)

fig.tight_layout()

```

```
[ ]: visualise_q(op_SMDP_q1[0],op_SMDP_1[0])
visualise_q(op_SMDP_q2[0],op_SMDP_2[0])
```

```
[ ]: visualise_q(op_SMDP_q1[1],op_SMDP_1[1])
visualise_q(op_SMDP_q2[1],op_SMDP_2[1])
```

```
[ ]: visualise_q(op_SMDP_q1[2],op_SMDP_1[2])
visualise_q(op_SMDP_q2[2],op_SMDP_2[2])
```

```
[ ]: visualise_q(op_SMDP_q1[3],op_SMDP_1[3])
visualise_q(op_SMDP_q2[3],op_SMDP_2[3])
op_SMDP_1[3]
```

```
[ ]: visualise_q(op_intra_q1[0],op_intra_1[0])
visualise_q(op_intra_q2[0],op_intra_2[0])
```

```
[ ]: visualise_q(op_intra_q1[1],op_intra_1[1])
visualise_q(op_intra_q2[1],op_intra_2[1])
```

```
[ ]: visualise_q(op_intra_q1[2],op_intra_1[2])
      visualise_q(op_intra_q2[2],op_intra_2[2])
```

```
[ ]: visualise_q(op_intra_q1[3],op_intra_1[3])
      visualise_q(op_intra_q2[3],op_intra_2[3])
```

```
[ ]: !sudo apt-get install texlive-xetex texlive-fonts-recommended_
      ↪texlive-plain-generic
```

```
[2]: # Run this only if you are using Google Colab
      from google.colab import drive
      import os

      drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: !jupyter nbconvert --to pdf /content/drive/MyDrive/Documents/Sem6-drive/RL/
      ↪Assignments/3Assignment/HRL1_final.ipynb
```

# HRL2\_final

April 23, 2022

[ ]:

[ ]:

```
[ ]: import numpy as np
import random
import gym
from gym.wrappers import Monitor
import glob
import io
import matplotlib.pyplot as plt
from IPython.display import HTML
import pandas as pd
from IPython import display
from IPython.display import clear_output
from tqdm import tqdm
from mpl_toolkits.axes_grid1 import make_axes_locatable
import matplotlib.patches as patches
```

```
[ ]: #Setting up the environment
env = gym.make('Taxi-v3')

env.reset()

#Current State
print(env.s)

# 4x12 grid = 48 states
print ("Number of states:", env.nS)

# Primitive Actions
action = ["south", "north", "east", "west", "pick", "drop"]
#correspond to [0,1,2,3] that's actually passed to the environment

# either go left, up, down or right
print ("Number of actions that an agent can take:", env.nA)
```

```

# Example Transitions
rnd_action = random.randint(0, 5)
print ("Action taken:", action[rnd_action])
next_state, reward, is_terminal, t_prob = env.step(rnd_action)
print ("Transition probability:", t_prob)
print ("Next state:", next_state)
print ("Reward recieved:", reward)
print ("Terminal state:", is_terminal)
env.render()
list(env.decode(env.s))

```

```

162
Number of states: 500
Number of actions that an agent can take: 6
Action taken: north
Transition probability: {'prob': 1.0}
Next state: 62
Reward recieved: -1
Terminal state: False

```

```

+-----+
|R: | : |G|
| : | : : |
| : : : : |
| | : | : |
|Y| : |B: |
+-----+
      (North)

```

```
[ ]: [0, 3, 0, 2]
```

```

[ ]: #0- "south"
      #1- "north"
      #2- "east"
      #3- "west"

```


```

[ ]: from IPython.display import Image

h_mat = np.array([[0,0,0,0,0],
                  [0,0,0,0,0],
                  [10,10,10,10,10],
                  [1,1,1,1,1],
                  [1,1,1,1,1]])
Image('/content/option_go_to_hw.png')


```

```
[ ]:
```




```
[ ]: Image('/content/option_left_hw.png')
```

```
[ ]:
```



```
[ ]: Image('/content/option_right_hw.png')
```

```
[ ]:
```



```
[ ]: # We are defining four more options here
      # Option 1 go to highway
      # Option 2 move left on highway
      # Option 3 move right on highway

def goto_hw(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = h_mat[ds[0],ds[1]]

    if (ds[0] == 2):    #termination condition(if row==2)
```



```

        optdone = True

    return [optact,optdone]

def left_hw(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = 3

    if (ds == [2,0]):      #termination condition
        optdone = True

    return [optact,optdone]

def right_hw(env,state):

    ds = list(env.decode(state))[:2]
    optdone = False
    optact = 2

    if (ds == [2,3]):      #termination condition
        optdone = True

    return [optact,optdone]

#Now the new action space will contain
#Primitive Actions: ["south", "north", "east", "west","pick","drop"]
#Options: ["move_to_R","move_to_Y","move_to_G","move_to_B"]
#Total Actions :["south", "north", "east",␣
→"west","pick","drop","goto_hw","left_hw","right_hw"]
#Corresponding to [0,1,2,3,4,5,6,7,8]

```

```

[ ]: #epsilon-greedy action selection function
seed = 36
rg = np.random.RandomState(seed)

def egreedy_policy(q_values,state,epsilon):
    nstate = list(env.decode(state))[:2]
    if nstate[0]!=2:
        actions = [0,1,2,3,4,5,6]
    if nstate == [2,0]:
        actions = [0,1,2,3,4,5,8]
    if nstate == [2,1] or nstate ==[2,2] :
        actions = [0,1,2,3,4,5,7,8]

```

```

if nstate == [2,3] or nstate == [2,4] :
    actions = [0,1,2,3,4,5,7]

if rg.rand() < epsilon:
    return rg.choice(actions)
else:
    #max = np.max(q_values[state])
    #return rg.choice(np.where(q_values[state] == max)[0])
    return np.argmax(q_values[state])

```

```

[ ]: def get_q():
    q_values_SMDP = np.zeros((500,9))
    for i in range(500):
        pos = list(env.decode(i))[:2]
        if pos[0] != 2:
            q_values_SMDP[i,7] = -100
            q_values_SMDP[i,8] = -100

        if pos == [2,0]:
            q_values_SMDP[i,7] = -100
            q_values_SMDP[i,6] = -100

        if pos == [2,1] or pos == [2,2] :
            q_values_SMDP[i,6] = -100

        if pos == [2,3] or pos == [2,4] :
            q_values_SMDP[i,6] = -100
            q_values_SMDP[i,8] = -100

    return q_values_SMDP

```

```

[ ]: ##### SMDP Q-Learning
actions = [0,1,2,3,4,5,6,7,8]
# Add parameters you might need here
gamma = 0.9
alpha = 0.2
epsilon=0.01
# Iterate over 1000 episodes
def SMDP(gamma,alpha,epsilon):
    q_values_SMDP =get_q()
    ufd1 = np.zeros((500,9))#Update_Frequency Data structure
    Rewards = []
    for _ in range(10000):
        state = env.reset()
        done = False

        # While episode is not over

```

```

episode_reward = 0
while not done:

    # Choose action
    action = egreedy_policy(q_values_SMDP, state, epsilon)

    # Checking if primitive action
    if action < 6:
        # Perform regular Q-Learning update for state-action pair

        next_state, reward, done, _ = env.step(action)
        q_values_SMDP[state, action] += alpha*(reward + gamma*np.
→max([q_values_SMDP[next_state, action] for action in actions]) -
→q_values_SMDP[state, action])
        ufd1[state,action] += 1
        state = next_state
        episode_reward += reward

    # Checking if action chosen is an option
    reward_bar = 0
    if action == 6: # action => goto_hw

        initial_state = np.copy(state)
        optdone = False
        count=0
        while (optdone == False):

            optact, _ = goto_hw(env, state)
            #
            next_state, reward, done, _ = env.step(optact)

            _, optdone = goto_hw(env, next_state)

            reward_bar = reward_bar + (gamma**count)*reward
            count+=1
            state = next_state
            episode_reward += reward

            q_values_SMDP[initial_state, action] += alpha*(reward_bar +
→(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]) -
→q_values_SMDP[initial_state, action])
            ufd1[initial_state,action] += 1

    if action == 7: # action => left_hw option

        initial_state = np.copy(state)

```

```

optdone = False
count=0
while (optdone == False):

    optact,_ = left_hw(env,state)

    next_state, reward, done,_ = env.step(optact)

    _,optdone = left_hw(env,next_state)

    reward_bar = reward_bar + (gamma**count)*reward
    count+=1

    state = next_state
    episode_reward += reward

    q_values_SMDP[initial_state, action] += alpha*(reward_bar +
↪(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
↪- q_values_SMDP[initial_state, action])
    ufd1[initial_state,action] += 1

if action == 8: # action => right_hw option

    initial_state = np.copy(state)
    optdone = False
    count=0
    while (optdone == False):

        optact,_ = right_hw(env,state)
        #
        next_state, reward, done,_ = env.step(optact)

        _,optdone = right_hw(env,next_state)

        reward_bar = reward_bar + (gamma**count)*reward
        count+=1

        state = next_state
        episode_reward += reward

        q_values_SMDP[initial_state, action] += alpha*(reward_bar +
↪(gamma**count)*np.max([q_values_SMDP[state, action] for action in actions]))
↪- q_values_SMDP[initial_state, action])
        ufd1[initial_state,action] += 1


```

```
    Rewards.append(episode_reward)

    return q_values_SMDP, Rewards, ufd1
```

```
[ ]: q_values1, rewards1, ufd1 = SMDP(gamma, alpha, epsilon)
plt.figure(figsize=(9,6))
plt.plot(rewards1)
```

```
[ ]: [ <matplotlib.lines.Line2D at 0x7fafee82e450>]
```



```
[ ]: np.mean(rewards1[8000:])
```

```
[ ]: 7.4725
```

```
[ ]: #### Intra-Option Q-Learning

# Add parameters you might need here
gamma = 0.93
alpha = 0.2
epsilon=0.001
# Iterate over 1000 episodes
def intra(gamma, alpha, epsilon):
    Rewards = []
```

```

q_values_SMDP2 = get_q()
ufd2 = np.zeros((500,9))#Update_Frequency Data structure
for _ in range(10000):
    state = env.reset()
    done = False

    # While episode is not over
    episode_reward = 0
    while not done:

        # Choose action
        action = egreedy_policy(q_values_SMDP2, state, epsilon)

        # Checking if primitive action
        if action < 6:
            # Perform regular Q-Learning update for state-action pair

            next_state, reward, done, _ = env.step(action)
            q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
→max([q_values_SMDP2[next_state, action] for action in actions]) -
→q_values_SMDP2[state, action])
            ufd2[state,action] += 1
            episode_reward+=reward

            state = next_state

        # Checking if action chosen is an option

        if action == 6: # action => goto highway option

            optdone = False
            while (optdone == False) :

                optact, _ = goto_hw(env, state)
                next_state, reward, done, _ = env.step(optact)
                _, optdone = goto_hw(env, next_state)
                episode_reward+=reward

                q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
→max([q_values_SMDP2[next_state, action] for action in actions]) -
→q_values_SMDP2[state, optact])
                ufd2[state, optact] += 1

                if not optdone:
                    q_values_SMDP2[state, action] += alpha*(reward +
→gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
                    ufd2[state, action] += 1

```

```

        else:
            q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions]) -
↪q_values_SMDP2[state, action])
            ufd2[state,action] += 1

            state = next_state

    if action == 7: # action => move left of highway option

        optdone = False
        while (optdone == False) :

            optact,_ = left_hw(env,state)
            next_state, reward, done,_ = env.step(optact)
            _,optdone = left_hw(env,next_state)
            episode_reward+=reward

            q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions]) -
↪q_values_SMDP2[state, optact])
            ufd2[state,optact] += 1

            if not optdone:
                q_values_SMDP2[state, action] += alpha*(reward +
↪gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
                ufd2[state,action] += 1
            else:
                q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions]) -
↪q_values_SMDP2[state, action])
                ufd2[state,action] += 1

            state = next_state

    if action == 8: # action => move right of highway option

        optdone = False
        while (optdone == False) :

            optact,_ = right_hw(env,state)
            next_state, reward, done,_ = env.step(optact)

```

```

        _,optdone = right_hw(env,next_state)
        episode_reward+=reward

        q_values_SMDP2[state, optact] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions]) -
↪q_values_SMDP2[state, optact])
        ufd2[state,optact] += 1

        if not optdone:
            q_values_SMDP2[state, action] += alpha*(reward +
↪gamma*q_values_SMDP2[next_state, action] - q_values_SMDP2[state, action])
            ufd2[state,action] += 1
        else:
            q_values_SMDP2[state, action] += alpha*(reward + gamma*np.
↪max([q_values_SMDP2[next_state, action] for action in actions]) -
↪q_values_SMDP2[state, action])
            ufd2[state,action] += 1

        state = next_state

    Rewards.append(episode_reward)

    return q_values_SMDP2,Rewards,ufd2

```

```

[ ]: q_values2,rewards2,ufd2 = intra(gamma,alpha,epsilon)
plt.figure(figsize=(9,6))
plt.plot(rewards2)


```

```

[ ]: [ <matplotlib.lines.Line2D at 0x7fafee7b0890>]

```





```
[ ]: np.mean(rewards2[8000:])
```

```
[ ]: 7.949
```

```
[ ]: #case when passenger is at either R,G,Y,B
#To visualise lets break the task into 2 parts: pick-up and drop
op_SMDP_1 = [np.zeros((5,5)) for i in range(4)] #Matrix to hold the optimal
→actions/options in each state
op_SMDP_q1 = [np.zeros((5,5)) for i in range(4)]
ufd_SMDP_1 = [np.zeros((5,5)) for i in range(4)]
op_intra_1 = [np.zeros((5,5)) for i in range(4)] #cases where passengers are
→not in the taxi
op_intra_q1 = [np.zeros((5,5)) for i in range(4)]
ufd_intra_1 = [np.zeros((5,5)) for i in range(4)]

op_SMDP_2 = [np.zeros((5,5)) for i in range(4)] #Matrix to hold the optimal
→actions/options in each state
op_SMDP_q2 = [np.zeros((5,5)) for i in range(4)]
ufd_SMDP_2 = [np.zeros((5,5)) for i in range(4)]
op_intra_2 = [np.zeros((5,5)) for i in range(4)] #cases where passenger is in
→the taxi
op_intra_q2 = [np.zeros((5,5)) for i in range(4)]
ufd_intra_2 = [np.zeros((5,5)) for i in range(4)]
```

```

for i in range(500): #iterating over all states
    state = list(env.decode(i))
    if state[2] != 4:
        op_SMDP_1[state[2]][state[0],state[1]] = np.argmax(q_values1[i])
        op_SMDP_q1[state[2]][state[0],state[1]] = np.amax(q_values1[i])
        ufd_SMDP_1[state[2]][state[0],state[1]] = np.sum(ufd1[i])
        op_intra_1[state[2]][state[0],state[1]] = np.argmax(q_values2[i])
        op_intra_q1[state[2]][state[0],state[1]] = np.amax(q_values2[i])
        ufd_intra_1[state[2]][state[0],state[1]] = np.sum(ufd2[i])
    else:
        op_SMDP_2[state[3]][state[0],state[1]] = np.argmax(q_values1[i])
        op_SMDP_q2[state[3]][state[0],state[1]] = np.amax(q_values1[i])
        ufd_SMDP_2[state[3]][state[0],state[1]] = np.sum(ufd1[i])
        op_intra_2[state[3]][state[0],state[1]] = np.argmax(q_values2[i])
        op_intra_q2[state[3]][state[0],state[1]] = np.amax(q_values2[i])
        ufd_intra_2[state[3]][state[0],state[1]] = np.sum(ufd2[i])

```

```

[ ]: def visualise_q(m1,m2):
    fig, ax = plt.subplots(figsize = (10,10))
    im = ax.imshow(m1, extent=[0, 10, 0, 10])
    ax.grid(which='major', axis='both', linestyle='-', color='k', linewidth=2)
    divider = make_axes_locatable(ax)
    cax = divider.append_axes('right', size='5%', pad=0.05)
    fig.colorbar(im, cax=cax, orientation='vertical')
    def x_direct(a):
        if a in [0,1,4,5,6,7,8]:
            return 0
        elif a in [2]:
            return 1
        else:
            return -1
    def y_direct(a):
        if a in [2,3,4,5,6,7,8]:
            return 0
        elif a in [1]:
            return 1
        else:
            return -1
    policyx = np.vectorize(x_direct)(m2)
    policyy = np.vectorize(y_direct)(m2)
    idx = 2*np.indices((5,5))
    ax.quiver(idx[1][::-1].ravel()+ 1, idx[0][::-1].ravel()+1, policyx.ravel(),
    ↪policyy.ravel(), pivot="middle", color='red')

    for i in range(5):
        for j in range(5):

```

```

    if m2[i][j] == 4: #pick-rectangle
        #rect = patches.Rectangle((idx[1][::-1][i][j]+0.5, idx[0][::-1][i][j]+0.
↪5), 1, 1, linewidth=0.5, edgecolor='r', facecolor='r')
        #ax.add_patch(rect)
        ax.text(idx[1][::-1][i][j]+0.6, idx[0][::-1][i][j]+0.6 , 'Pick', color =_
↪'r', fontsize= 20)
    if m2[i][j] == 5: #drop-Circle
        #circ = patches.Circle((idx[1][::-1][i][j]+1, idx[0][::-1][i][j]+1), 0.
↪5, linewidth=0.5, edgecolor='r', facecolor='r')
        #ax.add_patch(circ)
        ax.text(idx[1][::-1][i][j]+0.6, idx[0][::-1][i][j]+0.6 , 'Drop', color =_
↪'r', fontsize= 20)
    if m2[i][j] == 6:
        ax.text(idx[1][::-1][i][j]+0.9, idx[0][::-1][i][j]+0.9 , 'H', color =_
↪'r', fontsize= 35)

    if m2[i][j] == 7:
        ax.text(idx[1][::-1][i][j]+0.9, idx[0][::-1][i][j]+0.9 , 'L', color =_
↪'r', fontsize= 35)

    if m2[i][j] == 8:
        ax.text(idx[1][::-1][i][j]+0.9, idx[0][::-1][i][j]+0.9 , 'R', color =_
↪'r', fontsize= 35)

fig.tight_layout()

```

```


[ ]: def freq_heat(m1):
    fig, ax = plt.subplots(figsize = (10,10))
    im = ax.imshow(m1, extent=[0, 10, 0, 10])
    ax.grid(which='major', axis='both', linestyle='-', color='k', linewidth=2)
    divider = make_axes_locatable(ax)
    cax = divider.append_axes('right', size='5%', pad=0.05)
    fig.colorbar(im, cax=cax, orientation='vertical')


```

```


[ ]: visualise_q(op_SMDP_q1[0],op_SMDP_1[0])
    visualise_q(op_SMDP_q2[0],op_SMDP_2[0])


```






```
[ ]: visualise_q(op_SMDP_q1[1],op_SMDP_1[1])
      visualise_q(op_SMDP_q2[1],op_SMDP_2[1])
```







```
[ ]: visualise_q(op_SMDP_q1[2],op_SMDP_1[2])
      visualise_q(op_SMDP_q2[2],op_SMDP_2[2])
```










```
[ ]: visualise_q(op_SMDP_q1[3],op_SMDP_1[3])
      visualise_q(op_SMDP_q2[3],op_SMDP_2[3])
      freq_heat(ufd_SMDP_1[3])
      freq_heat(ufd_SMDP_2[3])
```












```
[ ]: visualise_q(op_intra_q1[0],op_intra_1[0])  
      visualise_q(op_intra_q2[0],op_intra_2[0])
```







```
[ ]: visualise_q(op_intra_q1[1],op_intra_1[1])
      visualise_q(op_intra_q2[1],op_intra_2[1])
      freq_heat(ufd_intra_2[1])
```

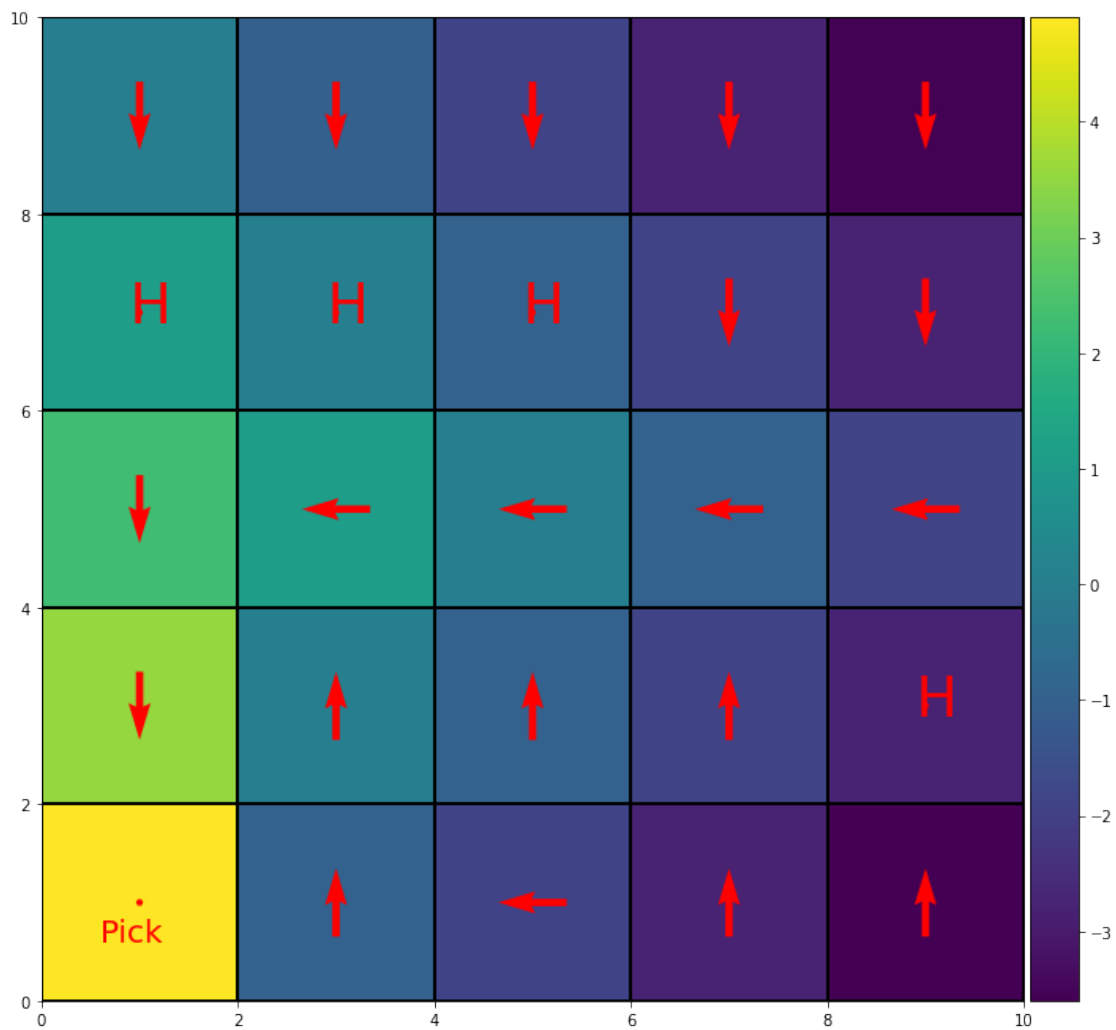










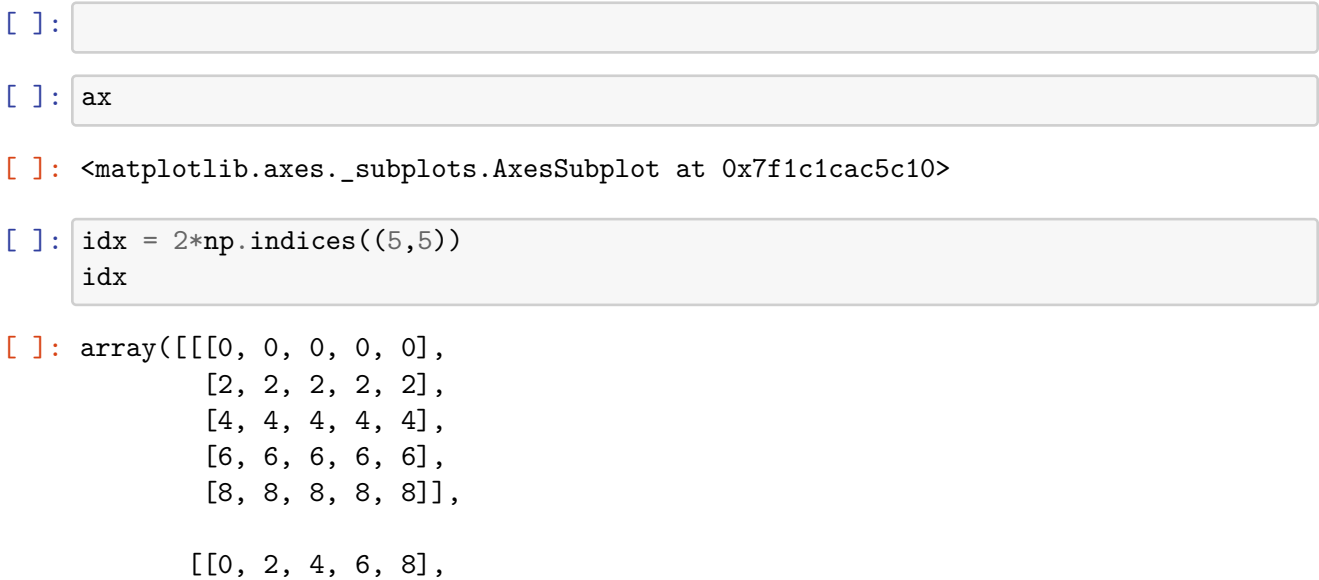
```
[ ]: visualise_q(op_intra_q1[2],op_intra_1[2])  
visualise_q(op_intra_q2[2],op_intra_2[2])
```





```
[ ]: visualise_q(op_intra_q1[3],op_intra_1[3])
      visualise_q(op_intra_q2[3],op_intra_2[3])
```





```
[0, 2, 4, 6, 8],  
[0, 2, 4, 6, 8],  
[0, 2, 4, 6, 8],  
[0, 2, 4, 6, 8]]])
```

## 1 Ignore

```
[ ]: !sudo apt-get install texlive-xetex texlive-fonts-recommended_  
↪texlive-plain-generic
```

```
[ ]: # Run this only if you are using Google Colab  
from google.colab import drive  
import os  
  
drive.mount('/content/drive')
```

```
[ ]: !jupyter nbconvert --to pdf /content/drive/MyDrive/Documents/Sem6-drive/RL/  
↪Assignments/3Assignment/HRL2_final.ipynb
```